

Proto-Personas — AI & Homework Integrity Project (v2)

Status: Proto-personas (assumption-based). Every claim is a hypothesis to be validated or killed in customer discovery interviews. Do not treat these as market facts yet.

Problem statement: High school teachers who assign out-of-class work struggle with students bypassing the cognitive effort required for learning, because AI allows them to submit work without real comprehension — and they want to ensure students are actually learning.

Market scope (team decision): Mexico first — good private preparatorias/bachilleratos (ages 15–18), not necessarily elite. Entry via Emiliano's network.

Design of the persona set (team decisions):

- **Teachers (3):** differentiated by *attitude toward AI* (resister / pragmatist / embracer) × *subject* (humanities / math-physics / English), with varied ages and personalities.
- **Students (3):** differentiated by *AI usage pattern* × *motivation* — heavy bypasser (grade optimizer), selective user (overloaded), honest low-user (anxious about false accusations).
- **Directors (3):** all Directors General, differentiated by *attitude toward AI policy* — prohibitionist / innovator / edtech-burned skeptic.

THE SCHOOLS

All nine personas live in one of three fictional schools (invented names, composite profiles — assumptions to validate like everything else here). Together they span the target segment: **good private preparatorias in mid-size Mexican cities** — solid, tuition-funded, academically respected, but not elite/international. Tuition in USD/month (MXN equivalent at ≈ 18.5, mid-2026).

Instituto Alborada — Beatriz's school (traditional / values-first)

- **Identity:** Founded 1962 by a religious congregation; K–12 on a single campus in an established, traditional neighborhood. Sells formation: discipline, values, character.
- **Size:** ~700 students total; ~180 in preparatoria; ~35 students per group.
- **Tuition:** USD **\$300–400/month** (≈ MXN 5,500–7,500) plus a modest annual inscription fee; conservative fee increases.
- **Families (NSE C+/B):** established middle class — shop and small-business owners, mid-level managers; many parents are alumni. They chose the school for order, safety, and values, not innovation.
- **Faculty & tech:** veteran, low-turnover faculty (Lourdes is an institution); deliberately low tech footprint — basic LMS, no 1:1 devices, phones banned in class.
- **Personas here:** Directora Beatriz, Profesora Lourdes, Rodrigo.

Colegio del Roble — Patricia's school (consolidated / balanced)

- **Identity:** Founded 1988, secular, UNAM-incorporated program; strong local academic reputation built over decades. Sells seriousness and stability.
- **Size:** ~900 students; ~260 in preparatoria; ~30 students per group.
- **Tuition:** USD **\$400–500/month** (≈ MXN 7,500–9,200); roughly 10% of prepa students hold partial merit scholarships (Valeria is one).
- **Families (NSE B):** professionals — doctors, engineers, accountants, established business owners. They expect academic rigor and measured, justified changes.

- **Faculty & tech:** stable veteran faculty that Patricia actively protects; some teachers (Andrés) also work at a second school. One well-adopted, teacher-chosen LMS + Google Workspace — and a graveyard of canceled platforms.
- **Personas here:** Directora Patricia, Profesor Andrés (plus a second, similar school across town), Valeria.

Lumen Bilingual Academy — Fernando's school (growth / tech-forward)

- **Identity:** Founded 2009 by an investor group; newer campus near expanding suburbs; marketing-heavy. Sells modernity: bilingual certification, tech electives, university placement.
 - **Size:** ~1,100 students and growing; ~320 in preparatoria; ~30 students per group.
 - **Tuition:** USD **\$500–650/month** (≈ MXN 9,200–12,000); premium tracks (bilingual certification, tech electives) sold on top.
 - **Families (NSE B/A-):** ambitious upper-middle class — dual-income professionals, entrepreneurs, many first-generation private-school parents highly sensitive to visible differentiation and placement results.
 - **Faculty & tech:** younger faculty, higher turnover; embracer teachers get visibility but little structural support. Google Workspace for Education, modern LMS, admissions CRM — a real budget for visible innovation.
 - **Personas here:** Director Fernando, Miss Renata, Santiago.
-

L

TEACHER · THE RESISTER (HUMANITIES)

Profesora Lourdes

Age	56
School	Instituto Alborada (Beatriz's)
Subject	Literatura y Historia de México
Experience	28 years teaching; deeply respected, somewhat feared
Load	4 groups × ~35 students
Personality	Principled, formal, proud of her craft; feels the profession is under attack
Tech comfort	Low — uses the school's LMS only because it's mandatory; grades on paper

CONTEXT

Lourdes built her career on the essay: reading a student's writing was how she *knew* them. Since 2023 that channel feels poisoned. She experiences AI not as a tool problem but as a **betrayal of the pact between teacher and student**. Her response has been to retreat: everything important now happens on paper, in class, under her eyes.

- GOALS**
- Restore trust: know with certainty that what she reads was thought by the student.
 - Protect the intellectual rigor of her subject — essays and deep reading must survive.
 - Finish her career with dignity, not as a "content moderator."
 - Have the school back her up with real, enforceable consequences.

- FRUSTRATIONS**
- She can *feel* AI-written text instantly but can **prove nothing** — and unprovable accusations have blown up with parents before.
 - In-class handwritten essays consumed the class time she used for discussion — her workaround is cannibalizing her actual teaching.
 - Younger colleagues who "embrace AI" strike her as surrendering; she feels isolated and old.
 - The school has no policy, so she's the "strict one" and students route their complaints around her.

- CURRENT BEHAVIOR / WORKAROUNDS**
- Moved ~80% of grading weight to handwritten, in-class work.
 - Refuses to grade anything typed at home; assigns less homework overall (and knows students read less as a result).
 - Occasionally interrogates suspected students orally — wins the battle, poisons the relationship.

- TECH & INFLUENCE**
- **Uses:** paper, whiteboard, mandated LMS (minimally), personal WhatsApp for parent groups.
 - **Adoption role:** *Vetoer*. She won't champion a tool, but her seniority means a tool she publicly rejects is dead in the sala de maestros. A tool she merely tolerates can survive.

"Yo no le tengo miedo a la inteligencia artificial. Le tengo miedo a graduar generaciones que nunca tuvieron una idea propia."

- ASSUMPTIONS TO VALIDATE**
1. Resisters are a large share of the teacher population (not a shrinking fringe) — this determines whether the product must win them or merely neutralize them.
 2. Retreat-to-paper is the dominant resister workaround, and teachers recognize its pedagogical cost.
 3. A resister would accept a tool framed as "restoring trust / proving authorship" even if she rejects anything framed as "integrating AI."

A

TEACHER · THE PRAGMATIST (MATH & PHYSICS)

Profesor Andrés

Age	41
School	Colegio del Roble (Patricia's) + a second, similar school across town
Subject	Matemáticas y Física
Experience	15 years; also tutors privately for extra income
Load	6 groups across two schools (~200 students)
Personality	Practical, tired, allergic to ideology; "just tell me what works"
Tech comfort	Moderate-high — Google Classroom, GeoGebra, Excel gradebook, has tried ChatGPT himself

CONTEXT

Photo-solver apps broke math homework *before* ChatGPT — Andrés has been fighting this war longer than the humanities teachers. Any problem set he sends home comes back perfect and means nothing. He's not angry anymore; he's **rebalancing**: homework barely counts, in-class exams count for almost everything. He knows this measures less (no long projects, no multi-step modeling) but it's honest.

- GOALS**
- Know which students can actually solve problems *before* the exam exposes them — early warning, not autopsy.
 - Keep homework pedagogically useful (practice matters in math) without pretending its grades mean anything.
 - Contain his workload: 200 students across two schools leaves zero slack for oral checks or process review.
 - Anything he adopts must save time net-net within two weeks or it's out.

- FRUSTRATIONS**
- Homework has become theater: he assigns it, AI solves it, he pretends to grade it.
 - The students who most need practice are exactly the ones outsourcing it — he watches them fail exams *predictably* and can't intervene earlier.
 - "Show your work" is dead: AI shows beautiful work.
 - Tools pitched at his school are always essay/plagiarism tools — nobody builds for math.

- CURRENT BEHAVIOR / WORKAROUNDS**
- Homework weight cut to 10%; exams and in-class quizzes carry the grade.
 - Randomizes numeric values per student on some worksheets (manual, painful, only slows students down for minutes).
 - Quick "explain step 3 to me" spot-checks when he smells outsourcing — effective, doesn't scale past a handful of students.

- TECH & INFLUENCE**
- **Uses:** Google Classroom, GeoGebra, Excel, WhatsApp, ChatGPT (for making exercises — an irony he acknowledges).
 - **Adoption role:** *Pragmatic influencer*. Won't fight for a tool, but if something demonstrably saves him time, he uses it and colleagues copy him. Fast to abandon anything with friction.

"La tarea ya no me dice quién sabe. Me dice quién tiene mejor prompt. Yo me entero de quién sabe hasta el examen — y ahí ya es demasiado tarde."

- ASSUMPTIONS TO VALIDATE**
1. STEM teachers' pain is *early visibility into who is actually practicing*, not authorship-proof (different product than for humanities).
 2. The "homework is now theater" rebalancing (weight → exams) is widespread in math/science.
 3. A net time-saving bar within ~2 weeks is the real adoption threshold for high-load teachers.
 4. Dual-school, ~200-student loads are common enough that "doesn't scale past 20 students" kills a feature.

Profesora Lourdes

 WHAT DO THEY THINK & FEEL?

- *What really counts*: students who can think for themselves; the craft and dignity of teaching.
- *Preoccupations*: she can no longer trust anything typed; every good essay is now a question mark.
- *Worries & aspirations*: graduating a generation without ideas of its own; finishing her career with dignity, with the essay intact.

 WHAT DO THEY HEAR?

- *Influencers*: media alarm about AI destroying education — it confirms what she already believes.
- *Boss*: Beatriz's AI-ban memo — a rule she agrees with but is given no means to enforce.
- *Friends/colleagues*: veteran peers trading suspicions; younger teachers saying "adapt or die," which she hears as an attack.

 WHAT DO THEY SEE?

- *Environment*: stacks of handwritten pages; class hours consumed by supervised writing instead of discussion.
- *Peers*: Renata-types embracing AI; other colleagues quietly giving up on homework altogether.
- *Market offers*: detector and "AI literacy" vendors — she distrusts both.

 WHAT DO THEY SAY & DO?

- *Toward others*: formal and demanding with students; more interrogative than she used to be (oral spot-checks that poison relationships).
- *In public*: the "strict one" in the sala de maestros; openly defends rigor, publicly rejects AI tools.
- *Appearance*: formal, traditional; paper gradebook under her arm — the analog teacher, on purpose.

 PAIN

- *Fears*: making an accusation she can't prove and having it blow up with parents; ending her career as a "content moderator."
- *Frustrations*: she can *feel* AI text instantly but can prove nothing; her workaround cannibalizes her own teaching time.
- *Obstacles*: no school policy with teeth, no evidence mechanism, growing isolation from younger colleagues.

 GAIN

- *Wants/needs*: certainty that what she reads was thought by the student; trust restored without becoming a police officer.
- *Measures of success*: essays she can grade at face value again; being respected rather than routed around.
- *Obstacles to overcome*: proving authorship without detectors, oral trials, or giving up take-home writing.

Profesor Andrés

 WHAT DO THEY THINK & FEEL?

- *What really counts*: knowing who can actually solve problems; honest measurement; protecting his scarce time.
- *Preoccupations*: 200 students across two schools; homework that has become theater he still has to stage.
- *Worries & aspirations*: watching students fail exams *predictably* without being able to intervene earlier; he wants early warning, not autopsy.

 WHAT DO THEY HEAR?

- *Influencers*: math-teacher WhatsApp/Facebook groups — everyone tells the same photo-solver war stories.
- *Boss*: two coordinators with two different (and contradictory) AI attitudes; neither offers him anything usable.
- *Friends/colleagues*: colleagues who quietly rebalanced everything to exams, same as he did.

 WHAT DO THEY SEE?

- *Environment*: perfect homework that means nothing; gradebooks full of numbers he doesn't believe.
- *Peers*: the whole department pretending to grade what AI pretended to solve.
- *Market offers*: essay/plagiarism tools pitched at his school — nothing built for math.

 WHAT DO THEY SAY & DO?

- *Toward others*: blunt, practical, fair; quick "explain step 3 to me" spot-checks when he smells outsourcing.
- *In public*: the no-ideology pragmatist — "just tell me what works"; complains little, adapts fast.
- *Appearance*: tired but reliable; the teacher who is always between two schools, laptop and Excel gradebook.

 PAIN

- *Fears*: that the students who most need practice keep outsourcing it and fail on schedule; that anything new eats his non-existent slack.
- *Frustrations*: "show your work" is dead — AI shows beautiful work; grading theater; manual value-randomizing that buys minutes.
- *Obstacles*: zero time for oral checks at 200 students; no tools designed for STEM; homework weight already cut to 10%.

 GAIN

- *Wants/needs*: early visibility into who is actually practicing, before the exam exposes them.
- *Measures of success*: net time saved within two weeks of adopting anything; fewer "predictable" exam failures.
- *Obstacles to overcome*: making homework a real signal again without adding a single hour to his week.

R

TEACHER · THE EMBRACER (ENGLISH)

Miss Renata

Age	29
School	Lumen Bilingual Academy (Fernando's)
Subject	English (B1–C1 groups; school sells bilingual certification to parents)
Experience	5 years; did a TESOL certificate, active in teacher communities online
Load	5 groups x ~30 students
Personality	Energetic, experimental, slightly contrarian; the "AI teacher" of her school
Tech comfort	High — builds her own materials with AI, runs a class Notion, early adopter of everything

CONTEXT

Renata thinks the war is over and AI won — so the assignment has to change, not the student. She's redesigned her homework around **process and performance**: drafts with visible revision, recorded voice notes, in-class conversations about what students wrote. Her problem isn't philosophy, it's **logistics**: her redesigned assessment doesn't scale, and the school gives her no time, no tools, and no cover when parents ask why her class "uses the cheating machine."

- GOALS
- Make AI-resistant (or AI-inclusive) assignments *sustainable* — right now her method costs her double the prep and grading time.
 - Prove to colleagues and dirección that her approach produces real language gains, with evidence.
 - Become the school's AI lead — she sees a career in it.
 - Give students a legitimate way to use AI (as tutor/feedback) without it doing the work for them.

- FRUSTRATIONS
- Reviewing process (drafts, revisions, recordings) for 150 students is crushing — she's the proof that the "right" pedagogy doesn't scale unaided.
 - Colleagues dismiss her as naive; Lourdes-types treat her as part of the problem.
 - No institutional policy means her AI-permissive rules contradict the AI-ban rules next door — students exploit the inconsistency.
 - Existing tools are detectors (which she thinks are snake oil) — nobody sells her what she actually needs: process visibility and scalable oral/formative assessment.

- CURRENT BEHAVIOR / WORKAROUNDS
- Requires Google Docs with version history; spot-checks revision timelines (tedious, gameable).
 - Voice-note assignments graded during her commute.
 - Wrote her own class AI policy ("use it to critique your draft, never to write it") — unenforceable, but sets culture.

- TECH & INFLUENCE
- **Uses:** Google Workspace, Notion, ChatGPT/Claude for materials, Vocaroo/WhatsApp voice notes, Kahoot.
 - **Adoption role:** *Champion*. She will pilot anything promising, evangelize it, and drag her school into a trial. The natural design partner and internal seller — but she has no budget authority.

“Prohibir la IA es como prohibir la calculadora en 1995. Mi problema no es que la usen — es que no tengo horas en el día para evaluar el proceso en vez del producto.”

- ASSUMPTIONS TO VALIDATE
1. Embracers exist in meaningful numbers in Mexican private schools and self-identify quickly (they're the beachhead if so).
 2. The embracer's bottleneck is genuinely *scaling process-based assessment*, not policy or conviction.
 3. A champion-teacher → director sales motion works (Renata can get a pilot approved).
 4. "Legitimate AI use" framing (tutor, not ghostwriter) resonates with parents rather than alarming them.

Miss Renata

 WHAT DO THEY THINK & FEEL?

- *What really counts*: real language gains; assessing process, not product; being on the right side of history.
- *Preoccupations*: drowning in drafts, revisions, and voice notes for 150 students — her "right" pedagogy doesn't scale.
- *Worries & aspirations*: burning out or being proven naive; becoming the school's (and eventually the region's) AI lead.

 WHAT DO THEY HEAR?

- *Influencers*: online TESOL and AI-in-education communities, edu-Twitter/TikTok — her real staffroom.
- *Boss*: Fernando showcases her at open house but gives her no hours, tools, or policy cover.
- *Friends/colleagues*: colleagues call her naive; Lourdes-types treat her as part of the problem.

 WHAT DO THEY SEE?

- *Environment*: students using AI everywhere regardless of rules; a school brochure promising what classrooms don't deliver.
- *Peers*: teachers next door banning what she permits — students exploiting the inconsistency.
- *Market offers*: detectors she considers snake oil; nobody selling process visibility or scalable oral assessment.

 WHAT DO THEY SAY & DO?

- *Toward others*: energetic, generous with students; gives them a legitimate way to use AI ("critique your draft, never write it").
- *In public*: the school's slightly contrarian "AI teacher"; evangelizes, pilots everything, presents at staff meetings.
- *Appearance*: young, digital-native, Notion-organized; grades voice notes on her commute.

 PAIN

- *Fears*: her method collapsing under its own weight; being stuck between parents who call AI "the cheating machine" and a school that won't back her.
- *Frustrations*: double prep and grading time; version-history spot-checks that are tedious and gameable; her class rules are unenforceable.
- *Obstacles*: no budget authority, no institutional policy, no tooling for process-based assessment at scale.

 GAIN

- *Wants/needs*: process-based, AI-resistant assessment that is *sustainable*; evidence her approach produces real gains.
- *Measures of success*: hours back in her week; measurable language outcomes she can show dirección; recognition as the school's AI lead.
- *Obstacles to overcome*: turning her one-classroom experiment into school-level practice — with data, not enthusiasm.

S

STUDENT · THE HEAVY BYPASSER (GRADE OPTIMIZER)

Santiago

Age	17, 5° semestre
School	Lumen Bilingual Academy (Fernando's)
Academic standing	Top-15% GPA — and it's inflated
Motivation	ITESM/UNAM admission; grades are currency, nothing else
AI usage	Daily, systematic; pays for ChatGPT Plus with allowance; has prompt templates per subject
Personality	Strategic, socially confident, sees himself as smart — "I optimize, I don't cheat"

CONTEXT

Santiago runs school like a portfolio. Every assignment gets a cost-benefit call: will effort here move my grade or my application? Essays, reports, reading reflections — AI, always, refined so it "sounds like him" (he keeps a prompt describing his own writing style). He does study for exams, because exams can't be outsourced — which to him proves the system: *what can be automated deserves to be.*

GOALS

- Maximum GPA per hour invested; protect time for extracurriculars that "look good."
- Never get caught — maintain plausible deniability (he paraphrases AI output enough to beat any detector).
- Get into a top university; nothing before that terminally matters.

FRUSTRATIONS

- Assignments that force synchronous effort (oral defenses, in-class writing, live problem-solving) — they break his model and he resents them as "gotchas."
- Teachers who vaguely accuse without proof — he's offended *performatively* and wins those exchanges.
- The quiet one: he sometimes bombs exam questions on material he "did homework on" and feels a flicker that his GPA is a bubble. He suppresses it.

CURRENT BEHAVIOR / WORKAROUNDS

- Prompt library per subject and teacher ("Lourdes flags fancy words — write at B2 level with two small grammar errors").
- Sells "homework help" (his prompts) to classmates informally.
- Studies genuinely only for exam-weighted material — a rational reading of Andrés-style grading policies.

TECH & INFLUENCE

- **Uses:** ChatGPT Plus, WhatsApp group chats (prompt-sharing economy), Google Docs, Photomath.
- **Adoption role:** *Adversary*. He will actively probe and route around any integrity tool; if he beats it, the whole class knows by Friday. A product that survives Santiago is credible.

“No es trampa, es eficiencia. La escuela califica el documento, no lo que aprendí. Yo entrego el documento.”

ASSUMPTIONS TO VALIDATE

1. Heavy bypassers are a minority, but their prompt-sharing sets the norms for everyone else (network effect of cheating).
2. Even hardened bypassers engage honestly when grading shifts to synchronous/oral formats — resentfully but compliantly.
3. The suppressed "my GPA is a bubble" anxiety is real and reachable (a wedge for products that make learning visible rather than punish).

V

STUDENT · THE SELECTIVE USER (OVERLOADED)

Valeria

Age	16, 3° semestre
School	Colegio del Roble (Patricia's), partial merit scholarship
Academic standing	Solid B+ student
Motivation	Survival of the load: 8 subjects, volleyball team, helps in the family business on weekends
AI usage	Selective — free ChatGPT for "busywork," honest effort on what feels real
Personality	Responsible, tired, a little cynical about school but not about learning

CONTEXT

Valeria's week has no slack. When homework piles up, something must give, and AI is the release valve — but she triages. The history questionnaire copied from a textbook? AI, zero guilt. Her chemistry lab write-up, in the subject she actually loves? Her own work, always. Her line isn't difficulty — it's **perceived pointlessness**. Assignments that feel like the teacher outsourced *their* effort get outsourced right back.

GOALS

- Keep grades good enough for a scholarship to stay at this school (real financial stakes).
- Protect time for the things that matter: chemistry, volleyball, family.
- Not fall behind in subjects she's automating — she knows there's a compounding risk and worries about it more than she admits.

FRUSTRATIONS

- Repetitive, copy-from-the-source assignments that insult her time — the direct trigger for AI use.
- Teachers who assign as if theirs is the only subject; no coordination of load across the 8.
- Feeling vaguely guilty *and* vaguely justified at once — she'd prefer a school where she didn't have to make this trade.
- Watching Santiago-types get better grades with zero effort while she rations her honesty.

CURRENT BEHAVIOR / WORKAROUNDS

- Triage system: honest effort on ~3 subjects she cares about or fears; AI-assisted or AI-done on the rest.
- Uses AI legitimately too — asks it to explain chemistry concepts her teacher rushed through.
- In her group chat, shares answers for "busywork" without hesitation; wouldn't share on "real" work.

TECH & INFLUENCE

- **Uses:** free ChatGPT, WhatsApp, Google Classroom, TikTok (where she actually learns study techniques).
- **Adoption role:** *Swing voter*. Valeria is the majority. A system that makes honest effort feel worthwhile converts her; a surveillance system just teaches her stealth. Product wins or dies on Valerias.

“Si la tarea es copiar del libro al cuaderno, la hace la IA y no me quita el sueño. Mi esfuerzo lo guardo para lo que sí me enseña algo.”

ASSUMPTIONS TO VALIDATE

1. Selective users are the largest student segment, and their triage line is *perceived assignment value*, not difficulty or subject.
2. Assignment redesign (make it feel worth doing) moves Valeria more than enforcement does.
3. Students already use AI legitimately (as tutor) and would trade transparency about it for protection from false accusations.
4. Load coordination failure across subjects is a root cause teachers/directors underestimate.

S

STUDENT · THE HEAVY BYPASSER (GRADE OPTIMIZER) · EMPATHY MAP

Santiago

 WHAT DO THEY THINK & FEEL?

- *What really counts*: admission to ITESM/UNAM; GPA-per-hour; his self-image as an optimizer, not a cheater.
- *Preoccupations*: running the cost-benefit call on every assignment; staying undetectable.
- *Worries & aspirations*: (suppressed) the flicker that his GPA is a bubble when he blanks on exam material "he did homework on"; the aspiration is one thing only — the admission letter.

 WHAT DO THEY HEAR?

- *Influencers*: admissions folklore and productivity/AI content on TikTok and YouTube; "work smart, not hard."
- *Boss/authority*: teachers' vague warnings about AI — he discounts them because they can never prove anything.
- *Friends*: prompt-sharing group chats where he is the acknowledged expert; classmates who pay for his "homework help."

 WHAT DO THEY SEE?

- *Environment*: a school that grades documents, not learning; assessment systems full of exploitable gaps.
- *Peers*: everyone using AI at different skill levels; the skilled ones thriving.
- *Market offers*: ever-better AI models and tools; detectors he has personally beaten.

 WHAT DO THEY SAY & DO?

- *Toward others*: charming and helpful — sells prompts, performs diligence to teachers, wins confrontations politely.
- *In public*: the model student — extracurriculars that "look good," offended (performatively) by unproven accusations.
- *Appearance*: polished, confident, curated — like his essays, tuned to "sound like him."

 PAIN

- *Fears*: getting caught with proof before applications close; synchronous formats (oral defenses, in-class writing) exposing the gap between his GPA and his knowledge.
- *Frustrations*: assignments that force real-time effort — he resents them as "gotchas."
- *Obstacles*: teachers shifting weight to in-class/exam work; any integrity system he can't game quickly.

 GAIN

- *Wants/needs*: maximum grades for minimum time, with plausible deniability intact.
- *Measures of success*: top GPA, the admission letter, peer status as the guy who beat the system.
- *Obstacles to overcome*: every new integrity mechanism — he will probe it, and if he beats it, the whole class knows by Friday.

V

STUDENT · THE SELECTIVE USER (OVERLOADED) · EMPATHY MAP

Valeria

 WHAT DO THEY THINK & FEEL?

- *What really counts*: keeping the scholarship; chemistry, volleyball, and her family — in that mental order when tired.
- *Preoccupations*: fitting 8 subjects, training, and the family business into a week with no slack.
- *Worries & aspirations*: the compounding gap in the subjects she automates (she worries more than she admits); a school where she wouldn't have to ration her honesty.

 WHAT DO THEY HEAR?

- *Influencers*: TikTok study-technique creators — where she actually learns how to learn.
- *Boss/authority*: eight teachers assigning as if theirs were the only subject; nobody coordinates the load.
- *Friends*: group chats where sharing answers for "busywork" is normal and unremarkable; parents counting on the scholarship.

 WHAT DO THEY SEE?

- *Environment*: copy-from-the-textbook assignments that insult her time; a stacked calendar.
- *Peers*: Santiago-types getting 10s with zero effort while she rations her honest hours.
- *Market offers*: free ChatGPT as the ever-present release valve; no tool that sees her workload.

 WHAT DO THEY SAY & DO?

- *Toward others*: responsible and helpful; shares busywork answers freely, never her "real" work.
- *In public*: the solid, dependable B+ student — nobody suspects how much she triages.
- *Appearance*: athlete's discipline, permanent slight tiredness; phone always at hand between obligations.

 PAIN

- *Fears*: losing the scholarship (real financial stakes); falling irrecoverably behind in automated subjects.
- *Frustrations*: pointless assignments — the direct trigger for AI use; the guilt-and-justification cocktail; watching effort go unrewarded.
- *Obstacles*: zero slack in the week; uncoordinated load across 8 subjects; a playing field tilted toward the cheaters.

 GAIN

- *Wants/needs*: assignments that feel worth doing; her honest effort to actually count for something.
- *Measures of success*: scholarship safe; learning in the subjects she loves; getting to sleep.
- *Obstacles to overcome*: the honest-effort penalty — a system where triage is the only rational strategy.

R

STUDENT · THE HONEST LOW-USER (ANXIOUS)

Rodrigo

Age	16, 4° semestre
School	Instituto Alborada (Beatriz's)
Academic standing	A- student, earned the hard way
Motivation	Genuinely likes learning (history, writing); wants to study law
AI usage	Minimal — occasional grammar check; won't let it write
Personality	Conscientious, quiet, rule-following; increasingly demoralized

CONTEXT

Rodrigo does the work — reads the chapters, writes his own essays, redoes problem sets until they click. He used to be near the top. Now his honest B+ essays sit next to AI-polished A essays, and teachers, newly suspicious of *good writing*, have twice questioned whether his work was his own. He was more insulted than scared — but also scared. The system punishes him from both directions: outcompeted by cheaters, suspected for being good.

GOALS

- Have his effort *count* — differentiation from students who outsource.
- Protection from false accusation: some way to prove his work is his, on his terms.
- Keep his intrinsic motivation alive in an environment that quietly tells him he's a sucker.

FRUSTRATIONS

- Grade compression: his honest work and Santiago's automated work are indistinguishable on paper — literally.
- Being asked "did you write this yourself?" for writing *well* — the accusation of quality.
- No channel to signal honesty: he can't opt in to scrutiny even though he'd welcome it.
- Watching his own resolve erode: "if I can't beat them and I get suspected anyway, why not join them?" He hasn't. Yet.

CURRENT BEHAVIOR / WORKAROUNDS

- Keeps handwritten notes and outline drafts as personal "evidence" in case he's accused again — self-invented process visibility.
- Started mentioning his drafts pre-emptively to teachers, which feels humiliating.
- Avoids AI even for legitimate uses (tutoring, feedback) because he fears any association.

TECH & INFLUENCE

- **Uses:** Google Docs, physical notebooks, WhatsApp (lurks in the answer-sharing chats, doesn't take).
- **Adoption role:** *Beneficiary and moral test*. Any tool that hurts Rodrigo (false flags, surveillance burden on the honest) is ethically and commercially broken — parents of Rodrigos are the ones who escalate. A tool he *loves* (lets him prove his work) has a built-in advocate base.

“Hice todo bien y me preguntaron si lo escribí yo. Los que lo hicieron con ChatGPT sacaron diez. Dime tú para qué me esfuerzo.”

ASSUMPTIONS TO VALIDATE

1. Honest students *want* a way to prove authorship (they'd opt into process visibility) rather than resenting it as surveillance.
2. False accusations of strong honest students are common enough to be a top-3 parent-escalation trigger.
3. Demoralization of honest students is a pain directors recognize and fear (it hits retention of exactly the families they want to keep).
4. The "erosion path" is real: honest students convert to bypassers over time when integrity isn't rewarded — making this urgent, not chronic.

Rodrigo

 WHAT DO THEY THINK & FEEL?

- *What really counts*: actually learning; his integrity; the law-school dream.
- *Preoccupations*: how to prove his work is his — before anyone asks again.
- *Worries & aspirations*: being formally accused and marked; his own eroding resolve ("why not join them?"); the aspiration that effort visibly counts.

 WHAT DO THEY HEAR?

- *Influencers*: teachers newly suspicious of good writing — the message that quality itself is now evidence of guilt.
- *Boss/authority*: a school (Alborada) that preaches honesty but can't distinguish his from Santiago-style output.
- *Friends*: classmates saying "just use it, everyone does"; parents telling him to keep his head down.

 WHAT DO THEY SEE?

- *Environment*: AI-polished 10s sitting next to his honest 8.5s; cheaters never caught.
- *Peers*: the answer-sharing chats he lurks in but never takes from.
- *Market offers*: nothing for him — every tool is aimed at catching cheaters, none at protecting the honest.

 WHAT DO THEY SAY & DO?

- *Toward others*: quiet, conscientious, rule-following; increasingly preemptive — mentions his drafts to teachers before being asked (and finds it humiliating).
- *In public*: keeps his head down; avoids AI even for legitimate uses, fearing any association.
- *Appearance*: handwritten notes and outline drafts he keeps as personal evidence — self-invented process visibility.

 PAIN

- *Fears*: a false accusation that sticks; losing the intrinsic motivation that defines him.
- *Frustrations*: the accusation of quality ("did you write this yourself?"); grade compression that makes honesty invisible.
- *Obstacles*: no channel to signal honesty or opt into scrutiny on his own terms.

 GAIN

- *Wants/needs*: a way to prove his work is his; differentiation from students who outsource.
- *Measures of success*: being trusted by default; grades that reflect effort; motivation intact through graduation.
- *Obstacles to overcome*: a system that punishes him from both directions — outcompeted by cheaters, suspected for being good.

B

DIRECTOR · THE PROHIBITIONIST

Directora Beatriz

Age	58
School	Instituto Alborada (see The Schools)
Role	Directora General, traditional private school (~700 students), strong religious/values identity
Accountable to	Owning family/board, parents who chose the school for discipline and values
Personality	Order-first, decisive, protective of the institution; technology is a managed risk, not an opportunity
Budget power	Full authority on purchases; conservative spender

CONTEXT

Beatriz's school sells formation: values, discipline, character. To her, AI-written homework is not a pedagogy question — it's **dishonesty**, full stop, and tolerating it corrodes the school's core promise. She issued a memo banning AI use on schoolwork. Enforcement, however, is delegated to teachers who have no way to enforce it, and she knows the memo is becoming a dead letter. Her nightmare is a public scandal in either direction: "school full of AI cheaters" or "school falsely accused my child."

- GOALS
- A rule that can actually be *enforced* — prohibition with teeth, applied uniformly and defensibly.
 - Zero scandals; keep AI conflicts out of parent WhatsApp groups and off social media.
 - Preserve the school's identity: rigor and honesty as the product parents pay for.
 - Legal safety: sanctions that survive an angry parent with a lawyer (minors' data, due process).

- FRUSTRATIONS
- Her ban is unenforceable and everyone quietly knows it — this undermines her authority, which she feels personally.
 - Detectors gave her teachers false positives; one parent meeting got ugly and now teachers under-report rather than risk it.
 - Vendors keep pitching her "AI literacy programs" — reading the room exactly backwards.
 - Teachers apply the policy inconsistently (Lourdes zealously, others not at all), which manufactures the parent complaints she fears most.

- CURRENT BEHAVIOR / WORKAROUNDS
- Written AI ban in the reglamento; sanctions defined but almost never applied (evidence problem).
 - Encouraged in-class, handwritten evaluation schoolwide — teachers complain about lost teaching time.
 - Handles each AI incident personally, case by case — unsustainable and inconsistent.

- TECH & INFLUENCE
- **School uses:** basic LMS, admin ERP; low tech footprint by design.
 - **Buying role:** *Economic buyer, values-gated*. Will pay for **enforceability, evidence, and defensibility** — never for "AI integration." Wins her with: due-process-proof evidence of authorship, fewer parent conflicts, uniform application. Loses her with: anything that puts AI in students' hands.

“En este colegio la honestidad no es negociable. Lo que necesito no es otro discurso sobre el futuro — necesito poder demostrar quién hizo la tarea.”

- ASSUMPTIONS TO VALIDATE
1. Prohibitionist directors are a substantial segment of good (non-elite) private schools — sized by values-identity schools.
 2. They will pay for *evidence/enforceability* even while rejecting AI-embracing framings (same product, different story?).
 3. The false-accusation legal fear is strong enough that "defensible evidence" beats "detection accuracy" as a message.
 4. A dead-letter ban is the common end-state of prohibition policies (validates that a ban creates demand rather than resolving it).

F

DIRECTOR · THE INNOVATOR

Director Fernando

Age	45
School	Lumen Bilingual Academy (see The Schools)
Role	Director General, growing private school (~1,100 students) competing hard for enrollment
Accountable to	Investor-owners focused on growth; ambitious middle-class parents
Personality	Marketer at heart; frames everything as differentiation; moves fast, sometimes faster than his faculty
Budget power	High, and actively looking to spend on visible innovation

CONTEXT

Fernando's school competes on being *modern*: bilingual, tech-forward, university-placement machine. He decided early that AI is an enrollment story — the school that "teaches AI responsibly" wins the open house. He's given talks to parents about it. The problem: behind the talk there is no operational substance. Teachers received one workshop and a one-page guideline; classrooms are as improvised as anywhere. He needs the reality to catch up with the brochure before parents notice the gap.

- GOALS
- A visible, parent-legible program: "here is how our school guarantees learning in the AI era" — with artifacts he can show at open house.
 - Enrollment growth and retention; differentiation from the traditional school across town (Beatriz's).
 - Teacher adoption that actually happens, so the story holds up when parents ask students.
 - Be first: he wants the case study, the press mention, the lighthouse status.

- FRUSTRATIONS
- The brochure-reality gap: he's promised "responsible AI education" and knows a skeptical parent could expose how thin it is.
 - Faculty drag: half his teachers quietly ban AI in their classrooms, contradicting the school's public position.
 - He can't measure anything: no data on actual AI use, actual learning, actual integrity — his story is all narrative, no evidence.
 - Vendors sell him either detectors (off-brand for him) or generic "AI literacy" content (no differentiation).

- CURRENT BEHAVIOR / WORKAROUNDS
- Published an "AI-responsible school" one-pager; runs parent talks; one teacher workshop per year.
 - Lets embracer teachers (his Renatas) run pilots informally and showcases them selectively.
 - Watches two competitor schools closely; fears one of them productizes this first.

- TECH & INFLUENCE
- **School uses:** Google Workspace for Education, a modern LMS, CRM for admissions — real tech budget exists.
 - **Buying role:** *Eager economic buyer*. Fastest sales cycle of the three — buys vision, needs proof later. Wins him with: parent-facing dashboards/certificates, "learning made visible" language, lighthouse-partner status. Risk: he'll buy shelf-ware if teachers don't adopt, then quietly churn (see Patricia).

“Los papás ya no preguntan si tenemos laboratorio de cómputo. Preguntan qué hacemos con la inteligencia artificial. El colegio que tenga la mejor respuesta se lleva las inscripciones.”

- ASSUMPTIONS TO VALIDATE
1. "Responsible AI" is genuinely an enrollment differentiator in Mexican private schools (parents choose partly on it) — or is it noise vs. price/placement/safety?
 2. Innovator directors will pay a premium for parent-facing evidence (dashboards, reports) layered on teacher tools.
 3. The brochure-reality gap creates urgency strong enough to close a sale in one budget cycle.
 4. Lighthouse/first-mover positioning ("founding school") accelerates the decision.

Directora Beatriz

WHAT DO THEY THINK & FEEL?

- *What really counts*: the institution's honor; honesty as non-negotiable; order and uniform rules.
- *Preoccupations*: her AI ban is becoming a dead letter — and everyone quietly knows it.
- *Worries & aspirations*: a public scandal in either direction ("school of cheaters" / "school falsely accused my child"); the aspiration is prohibition that actually works.

WHAT DO THEY HEAR?

- *Influencers*: fellow directors' scandal stories; diocesan/values-network circles reinforcing the dishonesty framing.
- *Boss*: the owning family/board, which expects discipline and zero reputational risk.
- *Friends/parents*: rumblings in parent WhatsApp groups — her early-warning system and her biggest fear.

WHAT DO THEY SEE?

- *Environment*: teachers applying her policy inconsistently — Lourdes zealously, others not at all — manufacturing the complaints she dreads.
- *Peers*: competitor schools (Lumen-types) marketing AI as a feature, which she reads as surrender.
- *Market offers*: detectors that burned her teachers with false positives; vendors pitching "AI literacy" — reading the room exactly backwards.

WHAT DO THEY SAY & DO?

- *Toward others*: decisive memos, personal case-by-case handling of every incident; protective of the school above all.
- *In public*: publicly firm — honesty is the product parents pay for; no visible doubt allowed.
- *Appearance*: formal, institutional authority; the reglamento is her instrument.

PAIN

- *Fears*: scandal on social media; an angry parent with a lawyer (minors' data, due process); her authority visibly hollowing out.
- *Frustrations*: an unenforceable rule; teachers under-reporting after the false-positive incident; tone-deaf vendors.
- *Obstacles*: no evidence mechanism that survives due process; inconsistent enforcement; a low-tech school by design.

GAIN

- *Wants/needs*: prohibition with teeth — defensible evidence of authorship, uniformly applied.
- *Measures of success*: zero AI conflicts escalating to parents or social media; sanctions that hold; the school's identity intact.
- *Obstacles to overcome*: enforcing honesty without triggering the very scandal she's preventing — and without putting AI in students' hands.

Director Fernando

WHAT DO THEY THINK & FEEL?

- *What really counts*: enrollment growth; differentiation; being first with the best answer to "what do you do with AI?"
- *Preoccupations*: the gap between the brochure and the classroom — he promised "responsible AI education" and knows how thin it is.
- *Worries & aspirations*: a skeptical parent (or a competitor) exposing the gap; the aspiration is lighthouse status — the case study, the press mention.

WHAT DO THEY SEE?

- *Environment*: his Renatas improvising pilots with no support; half his faculty quietly banning what the school publicly embraces.
- *Peers*: competitor schools circling the same differentiation story.
- *Market offers*: detectors (off-brand for him) and generic "AI literacy" content (no differentiation) — nothing he can showcase.

PAIN

- *Fears*: the brochure-reality gap exposed publicly; buying shelf-ware that teachers never adopt and churning quietly.
- *Frustrations*: faculty drag contradicting the public position; zero data behind his narrative — all story, no evidence.
- *Obstacles*: teacher adoption; no operational substance behind the marketing; a story students could contradict if asked.

WHAT DO THEY HEAR?

- *Influencers*: edtech conferences, marketing circles, what the two competitor schools are doing.
- *Boss*: investor-owners who read enrollment numbers, not pedagogy.
- *Friends/parents*: ambitious parents at open house asking about AI the way they used to ask about computer labs.

WHAT DO THEY SAY & DO?

- *Toward others*: moves fast, sometimes faster than his faculty; showcases pilot teachers selectively.
- *In public*: the visionary — parent talks, one-pagers, "AI-responsible school" branding.
- *Appearance*: polished, marketer's energy; the director who presents slides at open house himself.

GAIN

- *Wants/needs*: a visible, parent-legible program — artifacts, dashboards, certificates he can show at open house.
- *Measures of success*: enrollment up, retention up; a case study with his school's name on it; the story surviving when parents ask students.
- *Obstacles to overcome*: making reality catch up with the brochure before anyone checks.

P

DIRECTOR · THE EDETECH-BURNED SKEPTIC

Directora Patricia

Age	50
School	Colegio del Roble (see The Schools)
Role	Directora General, well-regarded mid-size private school (~900 students); 12 years in the role
Accountable to	A board of former parents; a faculty she's built and protects
Personality	Measured, evidence-driven, loyal to teachers; has seen every edtech wave come and go
Budget power	Full, but every purchase carries scar tissue

CONTEXT

Patricia has bought it all: the adaptive-learning platform (abandoned in one semester), the anti-plagiarism suite (false positives, canceled), the tablets program (drawer-ware). Each time, the pattern was identical — great demo, excited pilot teacher, faculty-wide rollout, silent death. She absolutely recognizes the AI-homework problem; her coordinadores raise it monthly and she's watched honest students' parents complain about unfair grading. But her operating principle is now: **no tool solves a practice problem**. Whatever she buys next must prove teachers will still be using it in month six.

GOALS

- Solve the real problem — she wants learning integrity, not a purchased alibi — but only via something her faculty will genuinely sustain.
- Protect teacher time and trust; she will not impose another platform on them top-down.
- Predictable, small-footprint spending: pilot first, pay for outcomes, easy exit.
- Keep the AI conflict from damaging her strongest asset: a stable, veteran faculty (her Lourdes and her Andrés).

FRUSTRATIONS

- Vendor fatigue: everyone's deck says "AI-powered"; nobody shows month-six retention data from a school like hers.
- She's caught between her own teachers: resisters demand bans, embracers demand tools, and any purchase takes a side.
- The problem is real and growing — doing nothing is also a decision, and she knows it's the wrong one.
- Detectors specifically burned her: a false positive nearly cost her a good family and a good teacher in the same week.

CURRENT BEHAVIOR / WORKAROUNDS

- Formed an internal AI committee (three teachers, includes an embracer and a resister) — slow, but builds legitimacy.
- Drafted an interim "teacher's discretion" policy — honest about being a placeholder.
- Takes vendor meetings but hasn't bought anything AI-related; asks every vendor for references from comparable Mexican schools and rarely gets them.

TECH & INFLUENCE

- **School uses:** LMS (well-adopted, chosen by teachers), Google Workspace; graveyard of canceled licenses.
- **Buying role:** *Hardest and most valuable buyer*. Slow cycle, pilot-gated, teacher-veto-empowered. Wins her with: bottom-up teacher demand (her own Renata asking for it), month-six usage data, references, per-classroom pricing, easy exit terms. A product that closes Patricia has proven adoption — she is the credibility beachhead for the whole market.

“Cada plataforma que he comprado me la vendieron como la solución definitiva. La única tecnología que sobrevive en mi colegio es la que los maestros piden por su nombre.”

ASSUMPTIONS TO VALIDATE

1. Edtech-burned skeptics are the modal director in good private schools (if so, adoption evidence is the product's core sales asset, not a nice-to-have).
2. Bottom-up teacher demand (champion asks director) genuinely unlocks her budget — validating a teacher-first go-to-market.
3. Pilot-priced, easy-exit offers materially shorten her cycle vs. annual-license offers.
4. Her teacher-veto dynamic means resister acceptance (Lourdes tolerating the tool) is a *purchase requirement*, not just a nice outcome.

Directora Patricia

 WHAT DO THEY THINK & FEEL?

- *What really counts*: her faculty's trust and time; real learning over purchased alibis; evidence over demos.
- *Preoccupations*: the AI problem is real and growing, but she cannot afford another failed platform.
- *Worries & aspirations*: another adaptive-learning-style corpse in the drawer; the aspiration is a tool her teachers ask for *by name* and still use in month six.

 WHAT DO THEY HEAR?

- *Influencers*: her coordinadores raising AI incidents monthly; the internal AI committee she formed.
- *Boss*: a board of former parents — patient, but watching how she handles this.
- *Friends/colleagues*: her own teachers pulling in opposite directions — resisters demanding bans, embracers demanding tools.

 WHAT DO THEY SEE?

- *Environment*: a graveyard of canceled licenses; the scar from a false positive that nearly cost a family and a teacher in one week.
- *Peers*: other directors buying shiny AI things she expects to see abandoned within a year.
- *Market offers*: every deck says "AI-powered"; nobody shows month-six retention data from a school like hers; references never materialize.

 WHAT DO THEY SAY & DO?

- *Toward others*: measured, loyal to teachers; refuses to impose platforms top-down.
- *In public*: takes every vendor meeting, buys nothing; asks each one for comparable Mexican-school references.
- *Appearance*: the steady, evidence-driven veteran — 12 years in the role and it shows.

 PAIN

- *Fears*: repeating the edtech cycle (demo → pilot → rollout → silent death); any purchase forcing her to take a side between her own teachers.
- *Frustrations*: vendor fatigue; the knowledge that doing nothing is also a decision — and the wrong one.
- *Obstacles*: scar tissue on every budget line; a pilot-gated process; a faculty with real veto power.

 GAIN

- *Wants/needs*: learning integrity genuinely solved — through something her faculty will sustain on their own.
- *Measures of success*: month-six usage data; teachers requesting the tool unprompted; easy exit terms she never has to use.
- *Obstacles to overcome*: buying again without breaking faculty trust — turning bottom-up demand into a safe purchase.

Synthesis

How the nine personas interlock

The classroom triangle (users x actors):

- **Lourdes** (resister) is in an arms race with **Santiago** (bypasser) — and her collateral damage is **Rodrigo** (falsely suspected). Her retreat-to-paper punishes **Valeria's** legitimate time constraints too.
- **Andrés** (pragmatist) has already conceded homework to Santiago and Valeria's triage; his pain is *not knowing who's practicing* until exams expose them.
- **Renata** (embracer) has the pedagogy that would convert Valeria, protect Rodrigo, and constrain Santiago — but it doesn't scale without tooling.

The buying triangle (buyers x everyone):

- **Beatriz** (prohibitionist) writes the policy Lourdes wants and can't enforce it; Santiago is her scandal risk, Rodrigo's parents her legal risk.
- **Fernando** (innovator) showcases Renata but hasn't equipped her; his brochure-reality gap is exposed the day a parent asks Valeria how school "teaches AI responsibly."
- **Patricia** (skeptic) will only buy what Renata demands *and* Lourdes tolerates — she is the adoption test the whole market imposes.

The strategic hypothesis (carried over, now sharpened): detection is a dead end — unreliable, adversarial versus Santiago, legally radioactive versus Rodrigo, already distrusted by Lourdes, Beatriz, and Patricia alike. The aligned frame is **making learning visible**: process evidence and assessment formats where honest effort is provable (Rodrigo), worthwhile (Valeria), scalable to grade (Renata, Andrés), enforceable (Lourdes, Beatriz), marketable (Fernando), and demonstrably adopted (Patricia). Every persona's incentives can align on that frame; none align on better catching.

Sharpest cross-persona tension to resolve in discovery: Beatriz will not buy anything that "embraces AI"; Fernando will not buy anything that smells like prohibition. Can one product serve both with different framing, or is this two segments and we must choose our beachhead?

Discovery interview plan (Round 1)

Channel: Emiliano's network → good private preparatorias in Mexico. Target ~20 conversations over 3–4 weeks.

SEGMENT	TARGET	SCREENING FOR	INTERVIEW SEED
Teachers	6–8 (aim: 2+ per attitude; cover humanities, STEM, English)	Assigns take-home work; 3+ years experience	Each teacher persona's "Assumptions to validate"
Students	5–6 (mix of semesters; recruit via teachers or siblings of network)	Uses or deliberately avoids AI on homework	Student personas' assumptions — especially Valeria's triage line and Rodrigo's opt-in appetite
Directors	3–4 (aim: at least one values-traditional, one growth-oriented)	Decision authority over edtech purchases	Director personas' assumptions — especially willingness-to-pay framing (evidence vs. innovation)

Kill criteria to watch for (evidence that should change our direction):

- Teachers consistently rank grading workload above integrity → the product is a grading tool with integrity as a feature, not the reverse.
- Students reject all process-visibility as surveillance (including Rodρίgos) → the "prove your work" wedge dies.
- Directors show no willingness to pay at school level (expect teachers to use free tools) → revisit segment or business model.
- The Beatriz/Fernando framing conflict proves unbridgeable → pick one beachhead segment before building.

Interview discipline: ask about past behavior ("tell me about the last time you suspected an AI-written assignment — what did you do?"), never about hypothetical futures ("would you use a tool that..."). Update or kill these personas after every 5 interviews.